

Drawing Inferences from Data:

Topics include:
statistical tests, pivot tables,
and Fisher's Exact test

Lecture 6

Statistical Hypothesis Tests

Test whether data support a give hypothesis or model.

For example:

"A £1 coin has equal probability of coming up heads or tails."



heads



tails

"A £1 coin has equal probability of coming up heads or tails."



heads



tails

By testing can we prove this statement to be true?

"A £1 coin has equal probability of coming up heads or tails."



heads



tails

By testing can we prove this statement to be true?

What if we do 10 trials and get 5 heads and 5 tails?

So, it looks like the statement is true.

"A £1 coin has equal probability of coming up heads or tails."



heads



tails

By testing can we prove this statement to be true?

What if we do 10 trials and get 5 heads and 5 tails?

So, it looks like the statement is true.

But, then we do 1,000,000 more trials and it turns out that the true probabilities are 55% heads and 45% tails.

In statistics, it is difficult to verify that a hypothesis is true.

"A £1 coin has equal probability of coming up heads or tails."



heads



tails

On the other hand, if the statement is false, can we show that?

If we were to flip this coin 1,000 times and get 763 heads and 237 tails, it's looking like the probabilities are not equal.

Some natural variation is expected, due to randomness of the individual flips. However, this is insufficient to explain the large difference between the heads and tails.

In statistics (and science), it would often be easier if hypotheses could be proven true.

However, we are often forced to work in the **reverse** direction, by **falsifying** hypothesis.

"In so far as a scientific statement speaks about reality, it must be falsifiable"

– Karl Popper

If a hypothesis is falsified with increasingly sensitive statistical tests, it *may* be true. (Or at least approximately true.)

Null Hypothesis Test

In statistics, we often want to determine whether something has an effect on data.

For example, in medicine, we may want to learn whether patients who get a particular treatment have a better outcome than patient who don't.

However, it may be difficult to directly test this, so we have to approach this problem from another direction.

Null Hypothesis Test

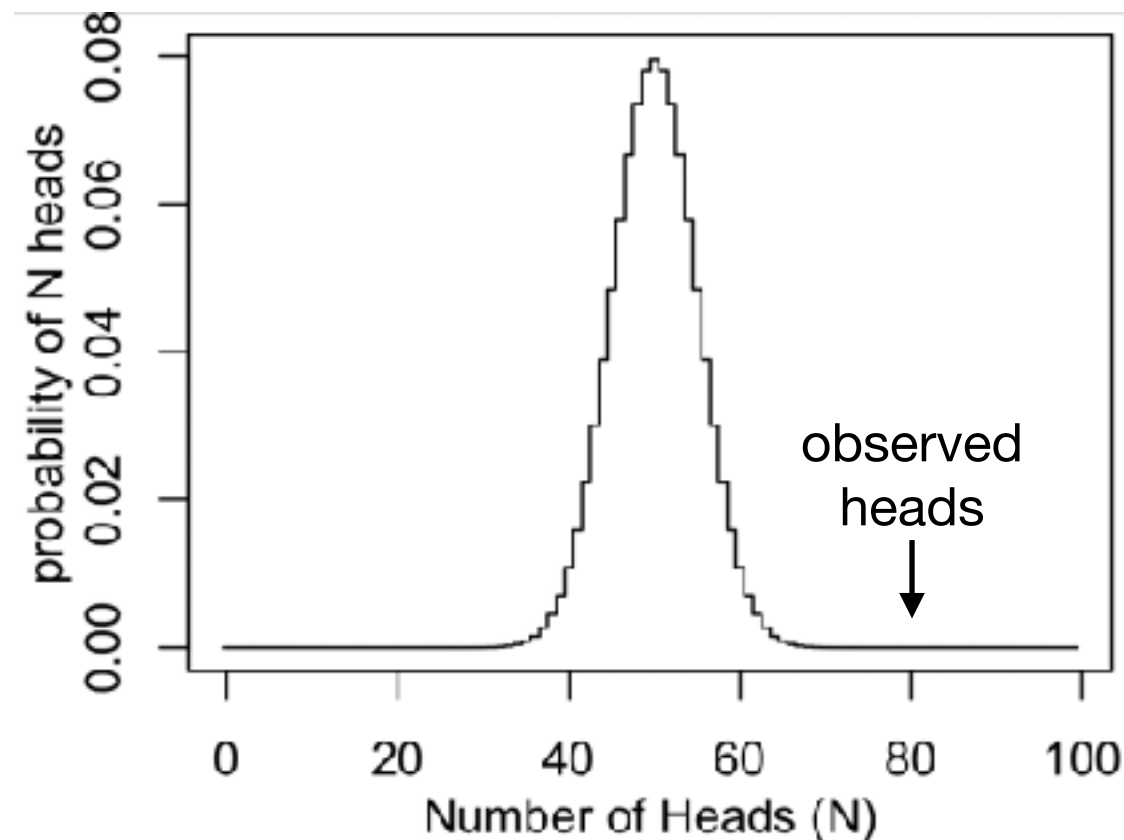
The "null hypothesis" is the hypothesis that there the treatment has not effect. (This is often denoted H_0 .)

The "alternative hypothesis" is that the treatment has an effect on the data. (Often denoted H_1 .)

We can more make a prediction ***assuming*** the null hypothesis. If there is a large mismatch between the prediction and the observed value (i.e., the probability of getting the observed value by chance is very small), we ***reject*** the null hypothesis, implying the alternative hypothesis is correct.

The coin example

Our null hypothesis is that a coin has a 50% probability of heads and a 50% probability of tails. Let's flip this coin 100 times. The probability of getting N heads is given by the following binomial distribution, where 50 is the peak of the distribution.



Suppose we observe 80 heads. Given the distribution above, the probability of observing 80 *or more* is $p=5 \times 10^{-10}$. This probability is called the ***p-value***.

This will effectively never happen by random chance, so, if we observed 80 heads, we can reject this

Fischer's Exact Test for Contingency Tables

		Test outcome		
		Passed	Failed	Row Total
Student preparation	Studied	45	10	55
	Did not study	9	10	19
	Column Total	54	20	74

Fischer's Exact Test for Contingency Tables

What is our null hypothesis?

There is no association between the two categorical variables. This implies that the probability of an observation falling into a particular category of one variable is independent of its classification in the other variable.

What is our alternate hypothesis?

There is an association between the two categorical variables.

We test this by assuming that the null hypothesis is true. Then, given this assumption, calculate the probability of obtaining the observed result.

If the probability of the total observed result or any more extremem result is very low, we can reject the null hypothesis. If the probability is not low, we say that our results are "consistent with the null hypothesis" and that no effect is detected.

Fischer's Exact Test for Contingency Tables

		Test outcome		
		Passed	Failed	Row Total
Student preparation	Studied	a	b	a+b
	Did not study	c	d	c+d
	Column Total	a+c	b+d	n=a+b+c+d

$$p = \frac{\binom{a+b}{a} \binom{c+d}{c}}{\binom{n}{a+c}} = \frac{\binom{a+b}{b} \binom{c+d}{d}}{\binom{n}{b+d}} = \frac{(a+b)! (c+d)! (a+c)! (b+d)!}{a! b! c! d! n!}$$

$$p_{\text{value}} = \sum_{\text{all tables as or more extreme}} p(\text{table})$$

		Test outcome		
		Passed	Failed	Row Total
Student preparation	Studied	a	b	a+b
	Did not study	c	d	c+d
Column Total		a+c	b+d	n=a+b+c+d

P values

Let's say we decide to use a probability threshold of 0.05 to reject the null hypothesis.

Which of the following would demonstrate a statistically significant effect?

$p=0.06$

$p=0.99$

$p=0.5$

$p=0.001$

$p=0.0001$

P values

Let's say we decide to use a probability threshold of 0.05 to reject the null hypothesis.

Which of the following would demonstrate a statistically significant effect?

$p=0.06$

$p=0.99$

$p=0.5$

$p=0.0001$

$p=0.001$

Fischer's Exact Test for Contingency Tables

		Test outcome		Row Total
		Passed	Failed	
Student preparation	Studied	45	10	55
	Did not study	9	10	19
Column Total		54	20	74

```
In [1]: import pandas as pd
....: from scipy.stats import fisher_exact
```

```
In [2]: data = {
....:     'Passed': [45, 9],
....:     'Failed': [10, 10]
....: }
```

```
In [3]: df = pd.DataFrame(data, index=['Studied', 'Did_Not_Study'])
```

```
In [4]: print(df)
           Passed  Failed
Studied         45      10
Did_Not_Study    9       10
```

```
In [5]: odds_ratio, p_value = fisher_exact(df.values)
....:
....: print(f"Odds Ratio: {odds_ratio:.4f}, P-value: {p_value:.4f}")
```

Odds Ratio: 5.0000, P-value: 0.0064



**And now for something
completely different**

Pivot Tables

Pandas Pivot Tables

Get mean value for *col3*,
indexed by *col1* and *col2*.

```
df.groupby(['col1', 'col2'])['col3'].aggregate('mean').unstack()
```

```
df.pivot_table('col3', index='col1', columns='col2')
```




The Titanic Dataset

```
In [1]: import numpy as np
...: import pandas as pd
...: import seaborn as sns
...: titanic = sns.load_dataset('titanic')
```

```
In [2]: titanic
```

```
Out[2]:
```

	survived	sex	age	fare	embarked	class	who	adult_male	embark_town	alive	alone
0	0	male	22.0	7.2500	S	Third	man	True	Southampton	no	False
1	1	female	38.0	71.2833	C	First	woman	False	Cherbourg	yes	False
2	1	female	26.0	7.9250	S	Third	woman	False	Southampton	yes	True
3	1	female	35.0	53.1000	S	First	woman	False	Southampton	yes	False
4	0	male	35.0	8.0500	S	Third	man	True	Southampton	no	True
...
886	0	male	27.0	13.0000	S	Second	man	True	Southampton	no	True
887	1	female	19.0	30.0000	S	First	woman	False	Southampton	yes	True
888	0	female	NaN	23.4500	S	Third	woman	False	Southampton	no	False
889	1	male	26.0	30.0000	C	First	man	True	Cherbourg	yes	True
890	0	male	32.0	7.7500	Q	Third	man	True	Queenstown	no	True

```
[891 rows x 11 columns]
```

I removed some columns for better display.

Comparing with GroupBy

```
In [8]: titanic.groupby(['sex', 'class'])['survived'].aggregate('mean').unstack()
```

```
.....:
```

```
Out[8]:
```

class	First	Second	Third
sex			
female	0.968085	0.921053	0.500000
male	0.368852	0.157407	0.135447

```
In [9]: titanic.pivot_table('survived', index='sex', columns='class')
```

```
.....:
```


```
Out[9]:
```

class	First	Second	Third
sex			
female	0.968085	0.921053	0.500000
male	0.368852	0.157407	0.135447

```
In [21]: titanic.pivot_table('survived', index=['sex'], columns=['class'],  
aggfunc='mean').plot.bar()
```

```
Out[21]: <AxesSubplot:xlabel='sex'>
```

```
In [22]: plt.show()
```



Multi-level Pivot Tables

```
In [23]: age = pd.cut(titanic['age'], [0, 18, 80]) # slices "age" column
      ...: age
```

Out[23]:

```
0      (18.0, 80.0]
1      (18.0, 80.0]
2      (18.0, 80.0]
3      (18.0, 80.0]
4      (18.0, 80.0]
```

```
      ...
886    (18.0, 80.0]
887    (18.0, 80.0]
888                NaN
889    (18.0, 80.0]
890    (18.0, 80.0]
```

Name: age, Length: 891, dtype: category

Categories (2, interval[int64, right]): [(0, 18] < (18, 80]]

```
In [24]: titanic.pivot_table('survived', ['sex', age], 'class')
```

Out[24]:

class		First	Second	Third
sex	age			
	female			
	(0, 18]	0.909091	1.000000	0.511628
	(18, 80]	0.972973	0.900000	0.423729
male	(0, 18]	0.800000	0.600000	0.215686
	(18, 80]	0.375000	0.071429	0.133663

Multi-level Pivot Tables

```
In [25]: titanic.pivot_table('alone', 'sex', ['embark_town', 'survived'])
```

```
Out[25]:
```

embark_town	Cherbourg		Queenstown		Southampton	
survived	0	1	0	1	0	1
sex						
female	0.222222	0.359375	0.666667	0.777778	0.301587	0.378571
male	0.696970	0.482759	0.736842	0.666667	0.750000	0.623377

Let's make a contingency table

Test whether 'sex' has an effect on 'survival'.

```
In [33]: test_tab = titanic.pivot_table(values= ?, index= ?,  
columns= ?, aggfunc= ? )
```

Let's make a contingency table

Test whether 'sex' has an effect on 'survival'.

```
In [33]: test_tab = titanic.pivot_table(values= ?, index='sex',  
columns= ?, aggfunc= ? )
```

Let's make a contingency table

Test whether 'sex' has an effect on 'survival'.

```
In [33]: test_tab = titanic.pivot_table(values= ?, index='sex',  
columns='survived', aggfunc= ? )
```

Let's make a contingency table

Test whether 'sex' has an effect on 'survival'.

```
In [33]: test_tab = titanic.pivot_table(values= ?, index='sex',  
columns='survived', aggfunc='count')
```

Let's make a contingency table

Test whether 'sex' has an effect on 'survival'.

```
In [33]: test_tab = titanic.pivot_table(values='class', index='sex',  
columns='survived', aggfunc='count')
```

```
In [34]: print(test_tab)
```

survived	0	1
sex		
female	81	233
male	468	109

Let's make a contingency table

Test whether 'sex' has an effect on 'survival'.

```
In [33]: test_tab = titanic.pivot_table(values='class', index='sex',  
columns='survived', aggfunc='count')
```

```
In [34]: print(test_tab)
```

survived	0	1
sex		
female	81	233
male	468	109

```
In [35]: odds_ratio, p_value = fisher_exact(test_tab.values)
```

```
In [36]: print(p_value)
```

6.463921564583144e-60

Let's make a contingency table

Test whether 'sex' has an effect on 'survival'.

```
In [33]: test_tab = titanic.pivot_table(values='class', index='sex',  
columns='survived', aggfunc='count')
```

```
In [34]: print(test_tab)  
survived    0    1  
sex  
female      81  233  
male       468  109
```

```
In [35]: odds_ratio, p_value = fisher_exact(test_tab.values)  
In [36]: print(p_value)  
6.463921564583144e-60
```

**Is this statistically
significant?**